

assignment_6.2b_cifar10

January 9, 2022

0.0.1 Assignment 6.2b

Using section 5.2 in Deep Learning with Python as a guide, create a ConvNet model that classifies images CIFAR10 small images classification dataset

- This time includes dropout and data-augmentation.
- Save the model, predictions, metrics, and validation plots in the dsc650/assignments/assignment06/results directory.

```
[1]: import json
      from pathlib import Path
      import os

      current_dir = Path(os.getcwd()).absolute()
      results_dir = current_dir.joinpath('results')

      print(current_dir)
      print(results_dir)
```

```
c:\Users\saman\git_repos\dsc650\dsc650\assignments\assignment06
c:\Users\saman\git_repos\dsc650\dsc650\assignments\assignment06\results
```

```
[2]: # loading the required libraries and packages
      import keras
      from keras.models import Sequential
      from keras.layers import Dense
      from keras.utils import to_categorical
      import matplotlib.pyplot as plt

      from keras.datasets import cifar10

      from keras.layers import Dropout
      from keras.layers import Flatten

      from keras.layers.convolutional import Conv2D
      from keras.layers.convolutional import MaxPooling2D

      from keras.utils import np_utils
      from keras.preprocessing.image import ImageDataGenerator
```

Using TensorFlow backend.

0.1 Data

Here we are using the CIFAR10 small images dataset to classify the images.

This is a dataset of 50,000 32X32 color training images and 10,000 test images labeled over 10 categories.

```
[3]: # Data preparation is required before training the model
def load_dataset():

    # loading the CIFAR10 dataset and create the training and test arrays
    (X_train, y_train), (X_test, y_test) = cifar10.load_data()

    # Lines 1 and 2 reshapes the inputs
    X_train = X_train.reshape((X_train.shape[0], 32, 32, 3)).
→astype('float32')
    X_test = X_test.reshape((X_test.shape[0], 32, 32, 3)).astype('float32')

    # Lines 3 and 4
    # Normalization of the input values (image pixels) from 0 and 255 to 0.1
    X_train = X_train / 255
    X_test = X_test / 255

    # Lines 5 and 6
    # one-hot encoding of the target variables
    y_train = np_utils.to_categorical(y_train)
    y_test = np_utils.to_categorical(y_test)

    num_classes = y_test.shape[1]

    return X_train, X_test, y_train, y_test
```

```
[8]: # loading the CIFAR10 dataset and create the training and test arrays
(X_train, y_train), (X_test, y_test) = cifar10.load_data()

# Lines 1 and 2 reshapes the inputs
X_train = X_train.reshape((X_train.shape[0], 32, 32, 3)).astype('float32')
X_test = X_test.reshape((X_test.shape[0], 32, 32, 3)).astype('float32')

# Lines 3 and 4
# Normalization of the input values (image pixels) from 0 and 255 to 0.1
X_train = X_train / 255
X_test = X_test / 255

# Lines 5 and 6
# one-hot encoding of the target variables
```

```

y_train = np_utils.to_categorical(y_train)
y_test = np_utils.to_categorical(y_test)

num_classes = y_test.shape[1]

```

```

[9]: # load_dataset()
print(f'Training set: {X_train.shape}')

```

Training set: (50000, 32, 32, 3)

```

[11]: def cnn_model_dropout():
    # function to create the CNN model
    # Create model
    model = Sequential() # model type is sequential
    # Stacking convolutional layers with small 3 X 3 filters
    # It is followed by a max pooling layer.
    # Each of the above blocks are repeated where the number of filters in
    → each block is increased.
    # Also the depth of the network such as 32,64 are also increased
    # Rectified Linear Activation ReLu is most widely used. It makes the
    → network sparse and efficient
    model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3),
    → activation='relu'))
    # model.add(Conv2D(32, (3, 3), activation='relu'))
    # Adding the pooling layer
    model.add(MaxPooling2D())
    model.add(Dropout(0.2))

    model.add(Conv2D(64, (3, 3), activation='relu'))
    # model.add(Conv2D(64, (3, 3), activation='relu'))
    # Adding the pooling layer
    model.add(MaxPooling2D())
    model.add(Dropout(0.2))

    model.add(Conv2D(128, (3, 3), activation='relu'))
    # model.add(Conv2D(128, (3, 3), activation='relu'))
    # Adding the pooling layer
    model.add(MaxPooling2D())
    model.add(Dropout(0.2))

    # Flatten layer converts the 2D matrix data to a vector
    model.add(Flatten())
    # Fully connected dense layer with 128 neurons
    model.add(Dense(128, activation='relu'))
    model.add(Dropout(0.2))
    # output layer which has 10 neurons for the 10 classes
    model.add(Dense(10, activation='softmax'))

```

```
return model
```

```
[21]: # Plotting the results
def summary_plot(history):

    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']
    loss = history.history['loss']
    val_loss = history.history['val_loss']

    epochs = range(1, len(acc) + 1)

    plt.plot(epochs, acc, 'bo', label='Training acc')
    plt.plot(epochs, val_acc, 'b', label='Validation acc')
    plt.title('Training and validation accuracy')
    plt.legend()

    plt.figure()

    plt.plot(epochs, loss, 'bo', label='Training loss')
    plt.plot(epochs, val_loss, 'b', label='Validation loss')
    plt.title('Training and validation loss')
    plt.legend()

    plt.show()
```

```
[5]: def summarize_diagnostics(history):
    plt.subplot(211)
    plt.title('Cross Entropy Loss')
    plt.plot(history.history['loss'], color='blue', label='train')
    plt.plot(history.history['val_loss'], color='orange', label='test')
    # plot accuracy
    plt.subplot(212)
    plt.title('Classification Accuracy')
    plt.plot(history.history['accuracy'], color='blue', label='train')
    plt.plot(history.history['val_accuracy'], color='orange', label='test')
    # save plot to file
    plt.savefig(f'{results_dir}\\2_plot.png')
    plt.show()
    plt.close()
```

```
[17]: def run_model():

    print('Load dataset')
    load_dataset()
```

```

print('dataset loaded')
print(f'Training set: {X_train.shape}')
print(f'Test Set: {X_test.shape}')
print(f'Number of categories : {num_classes}')

print('Build model')
# model = cnn_model()
model = cnn_model_dropout()
print('Model is defined')
print('Summary of the model.')
model.summary()

print('Compile Model')
model.compile(loss='categorical_crossentropy', optimizer='rmsprop',
↳metrics=['accuracy'])
print('Model compiled')

# Create data generator
datagen = ImageDataGenerator(
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest')
# prepare iterator
it_train = datagen.flow(X_train, y_train, batch_size=64)
steps = int(X_train.shape[0]/64)
print('Model fitting')
history = model.fit_generator(it_train, steps_per_epoch = steps,
↳validation_data=(X_test, y_test), epochs=15)

print('Saving the model')
model.save(f'{results_dir}\\assignment_6.2b_cifar10.h5')

print('Evaluating the model on the test data')
scores = model.evaluate(X_test, y_test, verbose=0)

print("CNN Accuracy: %.3f%%" % (scores[1]*100.0))

print('Output summary')
# summary_plot(history)
summarize_diagnostics(history)

```

[18]: run_model()

Load dataset

dataset loaded

Training set: (50000, 32, 32, 3)

Test Set: (10000, 32, 32, 3)

Number of categories : 10

Build model

Model is defined

Summary of the model.

Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d_13 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_10 (MaxPooling)	(None, 15, 15, 32)	0
dropout_11 (Dropout)	(None, 15, 15, 32)	0
conv2d_14 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_11 (MaxPooling)	(None, 6, 6, 64)	0
dropout_12 (Dropout)	(None, 6, 6, 64)	0
conv2d_15 (Conv2D)	(None, 4, 4, 128)	73856
max_pooling2d_12 (MaxPooling)	(None, 2, 2, 128)	0
dropout_13 (Dropout)	(None, 2, 2, 128)	0
flatten_3 (Flatten)	(None, 512)	0
dense_5 (Dense)	(None, 128)	65664
dropout_14 (Dropout)	(None, 128)	0
dense_6 (Dense)	(None, 10)	1290

Total params: 160,202

Trainable params: 160,202

Non-trainable params: 0

Compile Model

Model compiled

Model fitting

Epoch 1/15

781/781 [=====] - 60s 76ms/step - loss: 1.9537 -
accuracy: 0.2730 - val_loss: 1.6748 - val_accuracy: 0.3796

Epoch 2/15
781/781 [=====] - 59s 76ms/step - loss: 1.7110 - accuracy: 0.3750 - val_loss: 1.3704 - val_accuracy: 0.5188
Epoch 3/15
781/781 [=====] - 68s 88ms/step - loss: 1.6033 - accuracy: 0.4203 - val_loss: 1.3661 - val_accuracy: 0.5096
Epoch 4/15
781/781 [=====] - 65s 83ms/step - loss: 1.5326 - accuracy: 0.4498 - val_loss: 1.4312 - val_accuracy: 0.4861
Epoch 5/15
781/781 [=====] - 63s 80ms/step - loss: 1.4827 - accuracy: 0.4683 - val_loss: 1.1713 - val_accuracy: 0.5802
Epoch 6/15
781/781 [=====] - 67s 86ms/step - loss: 1.4497 - accuracy: 0.4844 - val_loss: 1.1538 - val_accuracy: 0.5875
Epoch 7/15
781/781 [=====] - 65s 83ms/step - loss: 1.4273 - accuracy: 0.4917 - val_loss: 1.1986 - val_accuracy: 0.5708
Epoch 8/15
781/781 [=====] - 64s 81ms/step - loss: 1.4080 - accuracy: 0.5003 - val_loss: 1.1596 - val_accuracy: 0.5884
Epoch 9/15
781/781 [=====] - 67s 86ms/step - loss: 1.3886 - accuracy: 0.5076 - val_loss: 1.1374 - val_accuracy: 0.5951
Epoch 10/15
781/781 [=====] - 64s 82ms/step - loss: 1.3890 - accuracy: 0.5099 - val_loss: 1.0909 - val_accuracy: 0.6222
Epoch 11/15
781/781 [=====] - 66s 85ms/step - loss: 1.3741 - accuracy: 0.5144 - val_loss: 1.0623 - val_accuracy: 0.6231
Epoch 12/15
781/781 [=====] - 65s 83ms/step - loss: 1.3646 - accuracy: 0.5205 - val_loss: 1.1935 - val_accuracy: 0.5840
Epoch 13/15
781/781 [=====] - 63s 81ms/step - loss: 1.3537 - accuracy: 0.5226 - val_loss: 1.0272 - val_accuracy: 0.6479
Epoch 14/15
781/781 [=====] - 76s 97ms/step - loss: 1.3585 - accuracy: 0.5218 - val_loss: 1.0183 - val_accuracy: 0.6468
Epoch 15/15
781/781 [=====] - 63s 80ms/step - loss: 1.3541 - accuracy: 0.5245 - val_loss: 1.0227 - val_accuracy: 0.6403
Saving the model
Evaluating the model on the test data
CNN Accuracy: 64.030%
Output summary

