

# assignment12

March 5, 2022

```
[1]: import keras
from keras.layers import Conv2D, Conv2DTranspose, Input, Flatten, Dense, Lambda, Reshape
from keras.models import Model
from keras.datasets import mnist
from keras import backend as K
import numpy as np
import matplotlib.pyplot as plt
```

Using TensorFlow backend.

```
[2]: # Load MNIST
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
[3]: # Normalize and Reshape
# Normalize

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train = x_train / 255
x_test = x_test / 255
```

```
[4]: # Reshape


img_width = x_train.shape[1]
img_height = x_train.shape[2]
num_channels = 1 # Grey scale data
x_train = x_train.reshape(x_train.shape[0], img_height, img_width, num_channels)
x_test = x_test.reshape(x_test.shape[0], img_height, img_width, num_channels)
input_shape = (img_height, img_width, num_channels)
```

```
[5]: #View a few images
plt.figure(1)
plt.subplot(331)
plt.imshow(x_train[54][:,:,0])

plt.subplot(332)
plt.imshow(x_train[555][:,:,0])
```

```
plt.subplot(333)
plt.imshow(x_train[6789][:,:,0])

plt.subplot(334)
plt.imshow(x_train[42000][:,:,0])
plt.show()
```



```
[6]: # Build the model
## Encoder
## Define 4 conv2D, flatten and ten dense
latent_dim = 2 # Number of latent dimension parms(latent space is a 2D plane)
input_img = Input(shape=input_shape, name='encoder_input')
x = Conv2D(32, 3, padding='same', activation='relu')(input_img)
x = Conv2D(64, 3, padding='same', activation='relu',strides=(2, 2))(x)
x = Conv2D(64, 3, padding='same', activation='relu')(x)
x = Conv2D(64, 3, padding='same', activation='relu')(x)
conv_shape = K.int_shape(x) # Shape of conv to be provided to decoder
# Flatten
x = Flatten()(x)
x = Dense(32, activation='relu')(x)
```

```
[7]: # Two outputs, for latent mean and log variance (std. dev.)
# Use these to sample random variables in latent space to which inputs are
↳ mapped.

z_mu = Dense(latent_dim, name='latent_mu')(x) # Mean values of encoded input
z_sigma = Dense(latent_dim, name='latent_sigma')(x) # Std dev. (variance) of
↳ encoded input
```

```
[8]: #REPARAMETERIZATION TRICK
# Define sampling function to sample from the distribution
# Reparameterize sample based on the process defined by Gunderson and Huang
# into the shape of: mu + sigma squared x eps
#This is to allow gradient descent to allow for gradient estimation accurately.
def sample_z(args):
    z_mu, z_sigma = args
    eps = K.random_normal(shape=(K.shape(z_mu)[0], K.int_shape(z_mu)[1]))
    return z_mu + K.exp(z_sigma / 2) * eps
```

```
[9]: # sample vector from the latent distribution
# z is the lambda custom layer we are adding for gradient descent calculations
# using mu and variance (sigma)
z = Lambda(sample_z, output_shape=(latent_dim, ), name='z')([z_mu, z_sigma])

#Z (lambda layer) will be the last layer in the encoder.
# Define and summarize encoder model.
encoder = Model(input_img, [z_mu, z_sigma, z], name='encoder')
print(encoder.summary())
```

Model: "encoder"

Layer (type)	Output Shape	Param #	Connected to
=====			
encoder_input (InputLayer)	(None, 28, 28, 1)	0	
-----			
conv2d_1 (Conv2D)	(None, 28, 28, 32)	320	
encoder_input[0][0]			
-----			
conv2d_2 (Conv2D)	(None, 14, 14, 64)	18496	conv2d_1[0][0]
-----			
conv2d_3 (Conv2D)	(None, 14, 14, 64)	36928	conv2d_2[0][0]
-----			
conv2d_4 (Conv2D)	(None, 14, 14, 64)	36928	conv2d_3[0][0]
-----			
flatten_1 (Flatten)	(None, 12544)	0	conv2d_4[0][0]
-----			
dense_1 (Dense)	(None, 32)	401440	flatten_1[0][0]
-----			
-----			

latent_mu (Dense)	(None, 2)	66	dense_1[0][0]
-----			
latent_sigma (Dense)	(None, 2)	66	dense_1[0][0]
-----			
z (Lambda)	(None, 2)	0	latent_mu[0][0]
latent_sigma[0][0]			
=====			
Total params: 494,244			
Trainable params: 494,244			
Non-trainable params: 0			
-----			
None			

```
[10]: # decoder takes the latent vector as input
decoder_input = Input(shape=(latent_dim,), name='decoder_input')
# Need to start with a shape that can be remapped to original image shape
# add dense layer with dimensions that can be resaped to desired output shape
x = Dense(conv_shape[1]*conv_shape[2]*conv_shape[3],
          activation='relu')(decoder_input)
# reshape to the shape of last conv. layer in the encoder, so we can
x = Reshape((conv_shape[1], conv_shape[2], conv_shape[3]))(x)
# upscale (conv2D transpose) back to original shape
# use Conv2DTranspose to reverse the conv layers defined in the encoder
x = Conv2DTranspose(32, 3, padding='same', activation='relu', strides=(2, 2))(x)
#Can add more conv2DTranspose layers, if desired.
#Using sigmoid activation
x = Conv2DTranspose(num_channels, 3, padding='same', activation='sigmoid',
                    name='decoder_output')(x)

# Define and summarize decoder model
decoder = Model(decoder_input, x, name='decoder')
decoder.summary()

# apply the decoder to the latent sample
z_decoded = decoder(z)
```

Model: "decoder"

Layer (type)	Output Shape	Param #
=====		
decoder_input (InputLayer)	(None, 2)	0
-----		
dense_2 (Dense)	(None, 12544)	37632
-----		

```

reshape_1 (Reshape)                (None, 14, 14, 64)                0
-----
conv2d_transpose_1 (Conv2DTr (None, 28, 28, 32)                18464
-----
decoder_output (Conv2DTransp (None, 28, 28, 1)                289
=====
Total params: 56,385
Trainable params: 56,385
Non-trainable params: 0
-----

```

```

[11]: #Define custom loss
#VAE is trained using two loss functions reconstruction loss and KL divergence
#Let us add a class to define a custom layer with loss
class CustomLayer(keras.layers.Layer):

    def vae_loss(self, x, z_decoded):
        x = K.flatten(x)
        z_decoded = K.flatten(z_decoded)

        # Reconstruction loss (as we used sigmoid activation we can use
        ↪binarycrossentropy)
        recon_loss = keras.metrics.binary_crossentropy(x, z_decoded)

        # KL divergence
        kl_loss = -5e-4 * K.mean(1 + z_sigma - K.square(z_mu) - K.exp(z_sigma),
        ↪axis=-1)
        return K.mean(recon_loss + kl_loss)

    # add custom loss to the class
    def call(self, inputs):
        x = inputs[0]
        z_decoded = inputs[1]
        loss = self.vae_loss(x, z_decoded)
        self.add_loss(loss, inputs=inputs)
        return x

```

```

[12]: # apply the custom loss to the input images and the decoded latent distribution
        ↪sample
y = CustomLayer()([input_img, z_decoded])
# y is basically the original image after encoding input img to mu, sigma, z
# and decoding sampled z values.
#This will be used as output for vae

```

```

[13]: vae = Model(input_img, y, name='vae')

# Compile VAE

```

```
vae.compile(optimizer='adam', loss=None)
vae.summary()
```

Model: "vae"

Layer (type)	Output Shape	Param #	Connected to
encoder_input (InputLayer)	(None, 28, 28, 1)	0	
conv2d_1 (Conv2D) encoder_input[0][0]	(None, 28, 28, 32)	320	
conv2d_2 (Conv2D)	(None, 14, 14, 64)	18496	conv2d_1[0][0]
conv2d_3 (Conv2D)	(None, 14, 14, 64)	36928	conv2d_2[0][0]
conv2d_4 (Conv2D)	(None, 14, 14, 64)	36928	conv2d_3[0][0]
flatten_1 (Flatten)	(None, 12544)	0	conv2d_4[0][0]
dense_1 (Dense)	(None, 32)	401440	flatten_1[0][0]
latent_mu (Dense)	(None, 2)	66	dense_1[0][0]
latent_sigma (Dense)	(None, 2)	66	dense_1[0][0]
z (Lambda) latent_sigma[0][0]	(None, 2)	0	latent_mu[0][0]
decoder (Model)	(None, 28, 28, 1)	56385	z[0][0]
custom_layer_1 (CustomLayer) encoder_input[0][0]	[(None, 28, 28, 1), 0		decoder[1][0]

=====

Total params: 550,629

Trainable params: 550,629

Non-trainable params: 0

-----  
-----

C:\Users\saman\.conda\envs\dsc650\lib\site-

packages\keras\engine\training\_utils.py:819: UserWarning: Output custom\_layer\_1 missing from loss dictionary. We assume this was done on purpose. The fit and evaluate APIs will not be expecting any data to be passed to custom\_layer\_1.

'be expecting any data to be passed to {0}.'.format(name))

```
[14]: # Train autoencoder
vae.fit(x_train, None, epochs = 10, batch_size = 32, validation_split = 0.2)
```

Train on 48000 samples, validate on 12000 samples

Epoch 1/10

48000/48000 [=====] - 109s 2ms/step - loss: 0.2315 -  
val\_loss: 0.2085

Epoch 2/10

48000/48000 [=====] - 118s 2ms/step - loss: 0.2021 -  
val\_loss: 0.1974

Epoch 3/10

48000/48000 [=====] - 105s 2ms/step - loss: 0.1953 -  
val\_loss: 0.1934

Epoch 4/10

48000/48000 [=====] - 112s 2ms/step - loss: 0.1919 -  
val\_loss: 0.1907

Epoch 5/10

48000/48000 [=====] - 110s 2ms/step - loss: 0.1895 -  
val\_loss: 0.1885

Epoch 6/10

48000/48000 [=====] - 117s 2ms/step - loss: 0.1880 -  
val\_loss: 0.1876

Epoch 7/10

48000/48000 [=====] - 113s 2ms/step - loss: 0.1866 -  
val\_loss: 0.1866

Epoch 8/10

48000/48000 [=====] - 108s 2ms/step - loss: 0.1856 -  
val\_loss: 0.1861

Epoch 9/10

48000/48000 [=====] - 113s 2ms/step - loss: 0.1847 -  
val\_loss: 0.1852


Epoch 10/10

48000/48000 [=====] - 120s 3ms/step - loss: 0.1839 -  
val\_loss: 0.1859

[14]: <keras.callbacks.callbacks.History at 0x1f556b73f98>


```
[15]: # =====  
# Visualize results  
# =====  
#Visualize inputs mapped to the Latent space  
#Remember that we have encoded inputs to latent space dimension = 2.  
#Extract z_mu --> first parameter in the result of encoder prediction  
#    ↳ representing mean  
  
mu, _, _ = encoder.predict(x_test)  
#Plot dim1 and dim2 for mu  
plt.figure(figsize=(10, 10))  
plt.scatter(mu[:, 0], mu[:, 1], c=y_test, cmap='brg')  
plt.xlabel('dim 1')  
plt.ylabel('dim 2')  
plt.colorbar()  
plt.show()
```





```
[16]: # Visualize images
#Single decoded image with random input latent vector (of size 1x2)
#Latent space range is about -5 to 5 so pick random values within this range
#Try starting with -1, 1 and slowly go up to -1.5,1.5 and see how it morphs
→from
#one image to the other.
sample_vector = np.array([[1,-1]])
decoded_example = decoder.predict(sample_vector)
decoded_example_reshaped = decoded_example.reshape(img_width, img_height)
plt.imshow(decoded_example_reshaped)
```

[16]: <matplotlib.image.AxesImage at 0x1f55f665da0>



```
[17]: #Let us automate this process by generating multiple images and plotting
      #Use decoder to generate images by tweaking latent variables from the latent_
      ↳space
      #Create a grid of defined size with zeros.
      #Take sample from some defined linear space. In this example range [-4, 4]
      #Feed it to the decoder and update zeros in the figure with output.
      import os
      from pathlib import Path

      n = 20 # generate 15x15 digits
      figure = np.zeros((img_width * n, img_height * n, num_channels))
      #Create a Grid of latent variables, to be provided as inputs to decoder.predict
      #Creating vectors within range -5 to 5 as that seems to be the range in latent_
      ↳space
      grid_x = np.linspace(-5, 5, n)
      grid_y = np.linspace(-5, 5, n)[::-1]

      # decoder for each square in the grid
      for i, yi in enumerate(grid_y):
          for j, xi in enumerate(grid_x):
              z_sample = np.array([[xi, yi]])
              x_decoded = decoder.predict(z_sample)
              digit = x_decoded[0].reshape(img_width, img_height, num_channels)
```

```

        figure[i * img_width: (i + 1) * img_width,
               j * img_height: (j + 1) * img_height] = digit

plt.figure(figsize=(15, 15))
#Reshape for visualization
fig_shape = np.shape(figure)
figure = figure.reshape((fig_shape[0], fig_shape[1]))


plt.imshow(figure, cmap='gnuplot2')

current_dir = Path(os.getcwd()).absolute()
results_dir = current_dir.joinpath('results')
results_dir.mkdir(parents=True, exist_ok=True)
vae_dir = results_dir.joinpath('vae')
vae_dir.mkdir(parents=True, exist_ok=True)

plt.savefig(f'{vae_dir}\\result.png')

plt.show()

```



[ ]: